

```

from datasets import load_dataset
from transformers import BertTokenizerFast, BertForSequenceClassification, Trainer, TrainingArguments
import torch
import numpy as np
import joblib
from sklearn.metrics import accuracy_score, f1_score, classification_report

```

```

MODEL_NAME = "bert-base-uncased"
NUM_LABELS = 4
MODEL_DIR = "bert_news_model"
DUMP_PATH = "news_classifier_metadata.joblib"
RANDOM_SEED = 42

```

```

# 1. Load dataset
print("Loading AG News dataset...")
dataset = load_dataset("ag_news")

```

↗ Loading AG News dataset...

```

# 2. Tokenizer
tokenizer = BertTokenizerFast.from_pretrained(MODEL_NAME)

def tokenize_batch(batch):
    return tokenizer(batch["text"], truncation=True, padding="max_length", max_length=128)

dataset = dataset.map(tokenize_batch, batched=True)

```

↗ Map: 100% 7600/7600 [00:01<00:00, 4841.95 examples/s]

```

# 3. Set format for PyTorch
dataset.set_format(type="torch", columns=["input_ids", "attention_mask", "label"])

```

```

train_dataset = dataset["train"]
test_dataset = dataset["test"]

```

```

# 4. Model
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)
model = BertForSequenceClassification.from_pretrained(MODEL_NAME, num_labels=NUM_LABELS)
model.to(device)

```

↗ Using device: cuda  
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```

BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSdpaSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
        )
      )
      (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
      )
      (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      )
    )
  )
)

```

```

        (dropout): Dropout(p=0.1, inplace=False)
    )
)
)
(pooler): BertPooler(
  (dense): Linear(in_features=768, out_features=768, bias=True)
  (activation): Tanh()
)
(dropout): Dropout(p=0.1, inplace=False)
(classifier): Linear(in_features=768, out_features=4, bias=True)
)

```

# 5. Metrics function

```

def compute_metrics(eval_pred):
    logits, labels = eval_pred
    preds = np.argmax(logits, axis=-1)
    acc = accuracy_score(labels, preds)
    f1 = f1_score(labels, preds, average="macro")
    return {"accuracy": acc, "f1_macro": f1}

```

```

training_args = TrainingArguments(
    output_dir="./results",
    eval_strategy="epoch",
    save_strategy="epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=32,
    num_train_epochs=3,
    weight_decay=0.01,
    load_best_model_at_end=True,
    seed=RANDOM_SEED,
    logging_dir="./logs",
    logging_steps=50,
    dataloader_num_workers=4,
    fp16=True
)

```

# 7. Trainer

```

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=test_dataset,
    processing_class=tokenizer,
    compute_metrics=compute_metrics
)

```

# 8. Train

```

trainer.train()

```



```
list-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes
i. (Learn how to deploy a W&B server locally: https://wandb.me/wandb-server)
I key in your browser here: https://wandb.ai/authorize?ref=models
m your profile and hit enter:wandb: WARNING If you're specifying your api key in code, ensure this code is no
ting the WANDB_API_KEY environment variable, or running `wandb login` from the command line.
creating one.
i.wandb.ai to your netrc file: /root/.netrc
as: moizpirzada11 (moizpirzada11-szabist) to https://api.wandb.ai. Use `wandb login --relogin` to force relog
0.21.0
tent/wandb/run-20250812_233742-b8i480rv
> Weights & Biases \(docs\)
moizpirzada11-szabist/huggingface
pirzada11-szabist/huggingface/runs/b8i480rv
list-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes
```

[15001/22500 32:06 < 16:03, 7.78 it/s, Epoch 2/3]

ation	Loss	Accuracy	F1 Macro
	0.177929	0.943289	0.943271
	0.191554	0.946316	0.946413

```
list-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes
```

```
list-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes
```

[22500/22500 48:50, Epoch 3/3]

ation	Loss	Accuracy	F1 Macro
	0.177929	0.943289	0.943271
	0.191554	0.946316	0.946413
	0.232124	0.946842	0.946889

```
00, training_loss=0.15091253214412265, metrics={'train_runtime': 3022.4575, 'train_samples_per_second':
econd': 7.444, 'total_flos': 2.368042020864e+16, 'train_loss': 0.15091253214412265, 'epoch': 3.0})
```

```
# 9. Evaluate
print("\n=== Final Evaluation on Test Set ===")
preds_output = trainer.predict(test_dataset)
y_true = preds_output.label_ids
y_pred = np.argmax(preds_output.predictions, axis=-1)
print("Accuracy:", accuracy_score(y_true, y_pred))
print("F1-macro:", f1_score(y_true, y_pred, average="macro"))
print(classification_report(y_true, y_pred, digits=4))
```



```
st Set ===
dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes

.06
132
recall  f1-score  support

0.9511   0.9543   1900
0.9905   0.9879   1900
0.9089   0.9140   1900
0.9226   0.9168   1900

          0.9433   7600
0.9433   0.9433   7600
0.9433   0.9433   7600
```


```
# 10. Save full model + tokenizer
model.save_pretrained(MODEL_DIR)
tokenizer.save_pretrained(MODEL_DIR)
print(f"Model + tokenizer saved to {MODEL_DIR}")
```



Model + tokenizer saved to bert\_news\_model

```
# 11. Save metadata with joblib
metadata = {
```

```
"model_dir": MODEL_DIR,  
"label_map": {0: "World", 1: "Sports", 2: "Business", 3: "Sci/Tech"},  
"base_model": MODEL_NAME  
}  
joblib.dump(metadata, DUMP_PATH)  
print(f"Metadata saved to {DUMP_PATH}")
```

 Metadata saved to news\_classifier\_metadata.joblib